

Лабораторная работа №2
по дисциплине
«Методы машинного обучения»
на тему
«Изучение библиотек обработки данных»

Выполнил:
студент группы ИУ5-24М
Лещев А. О.

1. Цель лабораторной работы

Изучить библиотеки обработки данных Pandas и PandaSQL [?].

2. Задание

Выполните первое демонстрационное задание “demo assignment” под названием “Exploratory data analysis with Pandas” со страницы купца <https://mlcourse.ai/assignments>

Условие задания - https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_engli

Официальный датасет находится здесь, но данные и заголовки хранятся отдельно, что неудобно для анализа - <https://archive.ics.uci.edu/ml/datasets/Adult>

Поэтому готовый набор данных для лабораторной работы удобнее скачать здесь - <https://raw.githubusercontent.com/Yorko/mlcourse.ai/master/data/adult.data.csv> (удобнее всего нажать на данной ссылке правую кнопку мыши и выбрать в контекстном меню пункт “сохранить ссылку”, будет предложено сохранить файл в формате CSV)

Пример решения задания - <https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution>

Набор упражнений по Pandas с решениями - https://github.com/guipsamora/pandas_exercises

3. Ход выполнения работы

Assignment #1 (demo)

Exploratory data analysis with Pandas

In this task you should use Pandas to answer a few questions about the Adult dataset.

Unique values of all features (for more information, please see the links above):

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.

- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- salary: >50K, <=50K.

Importing all required packages:

```
[1]: import pandas as pd
```

Setting maximum display width for text report [?]:

```
[2]: pd.set_option("display.width", 70)
```

Loading data:

```
[3]: data = pd.read_csv('data/adult.data.csv')
data.head()
```

```
[3]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	\
0	Never-married	Adm-clerical	Not-in-family	White	
1	Married-civ-spouse	Exec-managerial	Husband	White	
2	Divorced	Handlers-cleaners	Not-in-family	White	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	
4	Married-civ-spouse	Prof-specialty	Wife	Black	

	sex	capital-gain	capital-loss	hours-per-week	\
0	Male	2174	0	40	
1	Male	0	0	13	
2	Male	0	0	40	
3	Male	0	0	40	
4	Female	0	0	40	

	native-country	salary
0	United-States	<=50K
1	United-States	<=50K
2	United-States	<=50K
3	United-States	<=50K
4	Cuba	<=50K

1. How many men and women (sex feature) are represented in this dataset?

```
[4]: data["sex"].value_counts()
```

```
[4]: Male      21790
     Female    10771
     Name: sex, dtype: int64
```

2. What is the average age (age feature) of women?

```
[5]: data[data["sex"] == "Female"]["age"].mean()
```

```
[5]: 36.85823043357163
```

3. What is the percentage of German citizens (native-country feature)?

```
[6]: print("{0:%}".format(data[data["native-country"] == "Germany"]
                           .shape[0] / data.shape[0]))
```

```
0.420749%
```

4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature) and those who earn less than 50K per year?

```
[7]: ages1 = data[data["salary"] == "<=50K"]["age"]
     ages2 = data[data["salary"] == ">50K"]["age"]
     print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.
     →std()))
     print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.
     →std()))
```

```
<=50K: = 36.78373786407767 ± 14.02008849082488 years
>50K: = 44.24984058155847 ± 10.519027719851826 years
```

6. Is it true that people who earn more than 50K have at least high school education? (education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
[8]: high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",
                           "Assoc-voc", "Masters", "Doctorate"])
     def high_educated(e):
         return e in high_educations

     data[data["salary"] == ">50K"]["education"].map(high_educated).
     →all()
```

```
[8]: False
```

7. Display age statistics for each race (race feature) and each gender (sex feature). Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.

```
[9]: data.groupby(["race", "sex"])["age"].describe()
```

```
[9]:
```

	count	mean	std	min	\
race					
sex					

Amer-Indian-Eskimo	Female	119.0	37.117647	13.114991	17.0
	Male	192.0	37.208333	12.049563	17.0
Asian-Pac-Islander	Female	346.0	35.089595	12.300845	17.0
	Male	693.0	39.073593	12.883944	18.0
Black	Female	1555.0	37.854019	12.637197	17.0
	Male	1569.0	37.682600	12.882612	17.0
Other	Female	109.0	31.678899	11.631599	17.0
	Male	162.0	34.654321	11.355531	17.0
White	Female	8642.0	36.811618	14.329093	17.0
	Male	19174.0	39.652498	13.436029	17.0

		25%	50%	75%	max
race	sex				
Amer-Indian-Eskimo	Female	27.0	36.0	46.00	80.0
	Male	28.0	35.0	45.00	82.0
Asian-Pac-Islander	Female	25.0	33.0	43.75	75.0
	Male	29.0	37.0	46.00	90.0
Black	Female	28.0	37.0	46.00	90.0
	Male	27.0	36.0	46.00	90.0
Other	Female	23.0	29.0	39.00	74.0
	Male	26.0	32.0	42.00	77.0
White	Female	25.0	35.0	46.00	90.0
	Male	29.0	38.0	49.00	90.0

```
[10]: data[(data["race"] == "Amer-Indian-Eskimo")
          & (data["sex"] == "Male")]["age"].max()
```

```
[10]: 82
```

8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)? Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
[11]: def is_married(m):
        return m.startswith("Married")

data["married"] = data["marital-status"].map(is_married)
(data[(data["sex"] == "Male") & (data["salary"] == ">50K")])
["married"].value_counts()
```

```
[11]: True      5965
      False    697
      Name: married, dtype: int64
```

9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours, and what is the percentage of those who earn a lot (>50K) among them?

```
[12]: m = data["hours-per-week"].max()
      print("Maximum is {} hours/week.".format(m))

      people = data[data["hours-per-week"] == m]
```

```
c = people.shape[0]
print("{} people work this time at week.".format(c))

s = people[people["salary"] == ">50K"].shape[0]
print("{0:%} get >50K salary.".format(s / c))
```

Maximum is 99 hours/week.
 85 people work this time at week.
 29.411765% get >50K salary.

10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) for each country (native-country). What will these be for Japan?

```
[13]: p = pd.crosstab(data["native-country"], data["salary"],
                      values=data['hours-per-week'], aggfunc="mean")
p
```

```
[13]: salary
native-country
?          40.164760  45.547945
Cambodia   41.416667  40.000000
Canada     37.914634  45.641026
China      37.381818  38.900000
Columbia   38.684211  50.000000
Cuba       37.985714  42.440000
Dominican-Republic 42.338235  47.000000
Ecuador    38.041667  48.750000
El-Salvador 36.030928  45.000000
England    40.483333  44.533333
France     41.058824  50.750000
Germany    39.139785  44.977273
Greece     41.809524  50.625000
Guatemala  39.360656  36.666667
Haiti      36.325000  42.750000
Holand-Netherlands 40.000000      NaN
Honduras   34.333333  60.000000
Hong       39.142857  45.000000
Hungary    31.300000  50.000000
India      38.233333  46.475000
Iran       41.440000  47.500000
Ireland    40.947368  48.000000
Italy      39.625000  45.400000
Jamaica    38.239437  41.100000
Japan      41.000000  47.958333
Laos       40.375000  40.000000
Mexico     40.003279  46.575758
Nicaragua  36.093750  37.500000
Outlying-US(Guam-USVI-etc) 41.857143      NaN
Peru       35.068966  40.000000
Philippines 38.065693  43.032787
Poland     38.166667  39.000000
```

Portugal	41.939394	41.500000
Puerto-Rico	38.470588	39.416667
Scotland	39.444444	46.666667
South	40.156250	51.437500
Taiwan	33.774194	46.800000
Thailand	42.866667	58.333333
Trinidad&Tobago	37.058824	40.000000
United-States	38.799127	45.505369
Vietnam	37.193548	39.200000
Yugoslavia	41.600000	49.500000

```
[14]: p.loc["Japan"]
```

```
[14]: salary
      <=50K    41.000000
      >50K     47.958333
      Name: Japan, dtype: float64
```