

rk2

June 15, 2020

1 №2

, 5-24. №1.

1.1

CountVectorizer TfidfVectorizer.
(
LogisticRegression), Multinomial Naive Bayes (MNB), Complement Naive Bayes (CNB),
Bernoulli Naive Bayes.
(, accuracy).

1.2

1.2.1

```
[1]: from sklearn.datasets import fetch_20newsgroups
      from sklearn.feature_extraction.text import TfidfVectorizer

[2]: newsgroups_train = fetch_20newsgroups(subset='train', remove=('headers',
      ↳'footers'))
      newsgroups_test = fetch_20newsgroups(subset='test', remove=('headers',
      ↳'footers'))

[3]: vectorizer = TfidfVectorizer()
      vectorizer.fit(newsgroups_train.data + newsgroups_test.data)

[3]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
      dtype=<class 'numpy.float64'>, encoding='utf-8',
      input='content', lowercase=True, max_df=1.0, max_features=None,
      min_df=1, ngram_range=(1, 1), norm='l2', preprocessor=None,
```

```
smooth_idf=True, stop_words=None, strip_accents=None,
sublinear_tf=False, token_pattern='(?u)\\b\\w+\\b',
tokenizer=None, use_idf=True, vocabulary=None)
```

```
[4]: X_train = vectorizer.transform(newsgroups_train.data)
X_test = vectorizer.transform(newsgroups_test.data)

y_train = newsgroups_train.target
y_test = newsgroups_test.target
```

1.2.2

```
[5]: from sklearn.metrics import accuracy_score
```

```
[6]: def test(model):
    print(model)
    model.fit(X_train, y_train)
    print("accuracy:", accuracy_score(y_test, model.predict(X_test)))
```

```
[7]: from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB, ComplementNB, BernoulliNB
```

```
[8]: test(LogisticRegression(solver='lbfgs', multi_class='auto'))
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
accuracy: 0.774429102496017
```

```
[9]: test(MultinomialNB())
```

```
MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
accuracy: 0.72623473181094
```

```
[10]: test(ComplementNB())
```

```
ComplementNB(alpha=1.0, class_prior=None, fit_prior=True, norm=False)
accuracy: 0.8089484864577802
```

```
[11]: test(BernoulliNB())
```

```
BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)
accuracy: 0.5371747211895911
```

1.2.3

Complement Naive Bayes, ,
.